



Deploy the ASAv Auto Scale Solution on Microsoft Azure

- [Auto Scale Solution for ASAv on Azure, on page 1](#)
- [Download the Deployment Package, on page 3](#)
- [Auto Scale Solution Components, on page 3](#)
- [Auto Scale Solution Prerequisites, on page 4](#)
- [Auto Scale Deployment, on page 10](#)
- [Auto Scale Logic, on page 25](#)
- [Auto Scale Logging and Debugging, on page 25](#)
- [Auto Scale Guidelines and Limitations, on page 26](#)
- [Auto Scale Troubleshooting, on page 26](#)
- [Appendix - Build Azure Functions from Source Code, on page 27](#)

Auto Scale Solution for ASAv on Azure

About the Auto Scale Solution

ASAv Auto Scale for Azure is a complete serverless implementation which makes use of serverless infrastructure provided by Azure (Logic App, Azure Functions, Load Balancers, Security Groups, Virtual Machine Scale Set, etc.).

Some of the key features of the ASAv Auto Scale for Azure implementation include:

- Azure Resource Manager (ARM) template-based deployment.
- Support for scaling metrics based on CPU:



Note See [Auto Scale Logic, on page 25](#) for more information.

- Support for ASAv deployment and multi-availability zones.
- Completely automated configuration automatically applied to scaled-out ASAv instances.
- Support for Load Balancers and multi-availability zones.

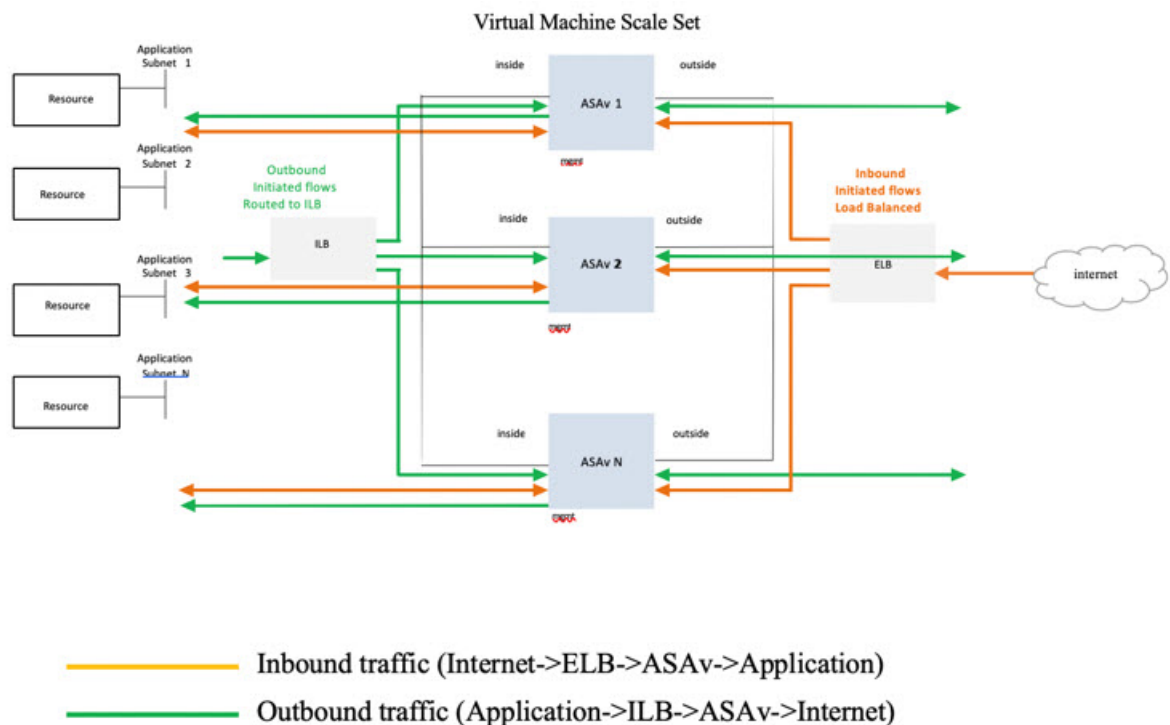
- Support for enabling and disabling the Auto Scale feature.
- Cisco provides an Auto Scale for Azure deployment package to facilitate the deployment.

Auto Scale Use Case

The ASAv Auto Scale for Azure is an automated horizontal scaling solution that positions an ASAv scale set sandwiched between an Azure Internal load balancer (ILB) and an Azure External load balancer (ELB).

- The ELB distributes traffic from the Internet to ASAv instances in the scale set; the firewall then forwards traffic to application.
- The ILB distributes outbound Internet traffic from an application to ASAv instances in the scale set; the firewall then forwards traffic to Internet.
- A network packet will never pass through both (internal & external) load balancers in a single connection.
- The number of ASAv instances in the scale set will be scaled and configured automatically based on load conditions.

Figure 1: ASAv Auto Scale Use Case



Scope

This document covers the detailed procedures to deploy the serverless components for the ASAv Auto Scale for Azure solution.

**Important**

- Read the entire document before you begin your deployment.
- Make sure the prerequisites are met before you start deployment.
- Make sure you follow the steps and order of execution as described herein.

Download the Deployment Package

The ASAv Auto Scale for Azure solution is an Azure Resource Manager (ARM) template-based deployment which makes use of the serverless infrastructure provided by Azure (Logic App, Azure Functions, Load Balancers, Virtual Machine Scale Set, etc.).

Download the files required to launch the ASAv Auto Scale for Azure solution. Deployment scripts and templates for your ASA version are available from the GitHub repository at:

- <https://github.com/CiscoDevNet/cisco-asav>

**Attention**

Note that Cisco-provided deployment scripts and templates for auto scale are provided as open source examples, and are not covered within the regular Cisco TAC support scope. Check GitHub regularly for updates and ReadMe instructions.

See [Appendix - Build Azure Functions from Source Code, on page 27](#) for instructions on how to build the *ASM_Function.zip* package.

Auto Scale Solution Components

The following components make up the ASAv Auto Scale for Azure solution.

Azure Functions (Function App)

The Function App is a set of Azure functions. The basic functionality includes:

- Communicate/Probe Azure metrics periodically.
- Monitor the ASAv load and trigger Scale In/Scale Out operations.

These functions are delivered in the form of compressed Zip package (see [Build the Azure Function App Package, on page 6](#)). The functions are as discrete as possible to carry out specific tasks, and can be upgraded as needed for enhancements and new release support.

Orchestrator (Logic App)

The Auto Scale Logic App is a workflow, i.e. a collection of steps in a sequence. Azure functions are independent entities and cannot communicate with each other. This orchestrator sequences the execution of these functions and exchanges information between them.

- The Logic App is used to orchestrate and pass information between the Auto Scale Azure functions.

- Each step represents an Auto Scale Azure function or built-in standard logic.
- The Logic App is delivered as a JSON file.
- The Logic App can be customized via the GUI or JSON file.

Virtual Machine Scale Set (VMSS)

The VMSS is a collection of homogeneous virtual machines, such as ASAv devices.

- The VMSS is capable of adding new identical VMs to the set.
- New VMs added to the VMSS are automatically attached with Load Balancers, Security Groups, and network interfaces.
- The VMSS has a built-in Auto Scale feature which is disabled for ASAv for Azure.
- You should not add or delete ASAv instances in the VMSS manually.

Azure Resource Manager (ARM) Template

The ARM template is used to deploy the resources required by the ASAv Auto Scale for Azure solution.

The ARM template provides input for the Auto Scale Manager components including:

- Azure Function App
- Azure Logic App
- The Virtual Machine Scale Set (VMSS)
- Internal/External load balancers.
- Security Groups and other miscellaneous components needed for deployment.



Important

The ARM template has limitations with respect to validating user input, hence it is your responsibility to validate input during deployment.

Auto Scale Solution Prerequisites

Azure Resources

Resource Group

An existing or newly created Resource Group is required to deploy all the components of this solution.



Note

Record the Resource Group name, the Region in which it is created, and the Azure Subscription ID for later use.

Networking

Make sure a virtual network is available/created. An Auto Scale deployment will not create, alter, or manage any networking resources.

The ASAv requires 3 network interfaces, thus your virtual network requires 3 subnets for:

1. Management traffic
2. Inside traffic
3. Outside traffic

The following ports should be open in the Network Security Group to which the subnets are connected:

- SSH(TCP/22)
Required for the Health probe between the Load Balancer and ASAv.
Required for communication between the Serverless functions and ASAv.
- Application-specific protocol/ports
Required for any user applications (for example, TCP/80, etc.).

**Note**

Record the virtual network name, the virtual network CIDR, the names of the 3 subnets, and the Gateway IP addresses of the outside and inside subnets.

Prepare the ASA Configuration File

Prepare an ASAv configuration file and store in a http/https server accessible by the ASAv instance. This is a standard ASA configuration file format. A scaled-out ASAv will download this file and update its configuration.

The ASA configuration file should have the following (at a minimum):

- Set DHCP IP assignment to all the interfaces.
- GigabitEthernet0/1 should be the 'inside' interface.
- GigabitEthernet0/0 should be the 'outside' interface.
- Set the gateway to the inside and outside interface.
- Enable SSH on the inside and outside interface from Azure utility IP (for health probe).
- Create a NAT configuration to forward traffic from the outside to the inside interface.
- Create an access policy to allow desired traffic.
- License the configuration. PAYG billing is not supported.

**Note**

There is no need to specifically configure the Management interface.

The following is a sample ASA configuration file for reference only.

```
ASA Version 9.13(1)
!
interface GigabitEthernet0/1
nameif inside
security-level 100
ip address dhcp setroute
!
interface GigabitEthernet0/0
nameif outside
security-level 0
ip address dhcp setroute
!
route outside 0.0.0.0 0.0.0.0 10.12.3.1 2
!
route inside 0.0.0.0 0.0.0.0 10.12.2.1 3
!
ssh 168.63.129.0 255.255.255.0 outside
!
ssh 168.63.129.0 255.255.255.0 inside
!
object network webserver
host 10.12.2.5
object service myport
service tcp source range 1 65535 destination range 1 65535
access-list outowebaccess extended permit object myport any any log disable
access-group outowebaccess in interface outside
object service app
service tcp source eq www
nat (inside,outside) source static webserver interface destination static interface any
service app app
object network obj-any
subnet 0.0.0.0 0.0.0.0
nat (inside,outside) source dynamic obj-any interface destination static obj-any obj-any
configure terminal
dns domain-lookup management
policy-map global_policy
class inspection_default
inspect icmp
call-home
profile License
destination transport-method http
destination address http https://tools.cisco.com/its/service/oddce/services/DDCEService
license smart
feature tier standard
throughput level 2G
license smart register idtoken <TOKEN>
: end
```

Build the Azure Function App Package

The ASAv Azure Auto Scale solution requires that you build an archive file: *ASM_Function.zip*, which delivers a set of discrete Azure functions in the form of compressed ZIP package.

See [Appendix - Build Azure Functions from Source Code, on page 27](#) for instructions on how to build the *ASM_Function.zip* package.

These functions are as discrete as possible to carry out specific tasks, and can be upgraded as needed for enhancements and new release support.

Input Parameters

The following table defines the template parameters and provides an example. Once you decide on these values, you can use these parameters to create the ASAv device when you deploy the ARM template into your Azure subscription. See [Deploy the Auto Scale ARM Template, on page 10](#).

Table 1: Template Parameters

Parameter Name	Allowed Values/Type	Description	Resource Creation Type
resourceNamePrefix	String* (between 3 and 10 characters)	All the resources are created with name containing this prefix. Note: Use only lowercase letters. Example: asav	New
virtualNetworkRg	String	The virtual network resource group name. Example: cisco-virtualnet-rg	Existing
virtualNetworkName	String	The virtual network name (already created). Example: cisco-virtualnet	Existing
mgmtSubnet	String	The management subnet name (already created). Example: cisco-mgmt-subnet	Existing
insideSubnet	String	The inside Subnet name (already created). Example: cisco-inside-subnet	Existing
internalLbIp	String	The internal load balancer IP address for the inside subnet (already created). Example: 1.2.3.4	Existing
outsideSubnet	String	The outside subnet name (already created). Example: cisco-outside-subnet	Existing
softwareVersion	String	The ASAv Version (selected from drop-down during deployment). Default: 914.1.0 Allowed: 914.1.0, 913.1.0	Existing
vmSize	String	Size of ASAv instance (selected from drop-down during deployment).	N/A

Parameter Name	Allowed Values/Type	Description	Resource Creation Type
asaAdminUserName	String*	<p>User name for the ASAv 'admin' user.</p> <p>Passwords must be 12 to 72 characters long, and must have: lowercase, uppercase, numbers, and special characters; and must have no more than 2 repeating characters.</p> <p>This cannot be 'admin'. See Azure for VM administrator user name guidelines.</p> <p>Note There is no compliance check for this in the template.</p>	New
asaAdminUserPassword	String*	<p>Password for the ASAv administrator user.</p> <p>Passwords must be 12 to 72 characters long, and must have: lowercase, uppercase, numbers, and special characters; and must have no more than 2 repeating characters.</p> <p>Note There is no compliance check for this in the template.</p>	New
scalingPolicy	POLICY-1 / POLICY-2	<p>POLICY-1: Scale-Out will be triggered when the average load of any ASAv goes beyond the Scale Out threshold for the configured duration.</p> <p>POLICY-2: Scale-Out will be triggered when average load of all the ASAv devices in the auto scale group goes beyond the Scale Out threshold for the configured duration.</p> <p>In both cases Scale-In logic remains the same: Scale-In will be triggered when average load of all the ASAv devices comes below the Scale In threshold for the configured duration.</p>	N/A

Parameter Name	Allowed Values/Type	Description	Resource Creation Type
scalingMetricsList	String	Metrics used in making the scaling decision. Allowed: CPU Default: CPU	N/A
scaleInThreshold	String	The Scale-In threshold in percent. Default: 10 When the ASAv metric goes below this value the Scale-In will be triggered. See Auto Scale Logic, on page 25 .	N/A
scaleOutThreshold	String	The Scale-Out threshold in percent. Default: 80 When the ASAv metric goes above this value, the Scale-Out will be triggered. The 'scaleOutThreshold' should always be greater than the 'scaleInThreshold'. See Auto Scale Logic, on page 25 .	N/A
minAsaCount	Integer	The minimum ASAv instances available in the scale set at any given time. Example: 2	N/A
maxAsaCount	Integer	The maximum ASAv instances allowed in the Scale set. Example: 10 Note The Auto Scale logic will not check the range of this variable, hence fill this carefully.	N/A

Parameter Name	Allowed Values/Type	Description	Resource Creation Type
metricsAverageDuration	Integer	<p>Select from the drop-down.</p> <p>This number represents the time (in minutes) over which the metrics are averaged out.</p> <p>If the value of this variable is 5 (i.e. 5min), when the Auto Scale Manager is scheduled it will check the past 5 minutes average of metrics and based on this it will make a scaling decision.</p> <p>Note Only numbers 1, 5, 15, and 30 are valid due to Azure limitations.</p>	N/A
initDeploymentMode	BULK / STEP	<p>Primarily applicable for the first deployment, or when the Scale Set does not contain any ASAv instances.</p> <p>BULK: The Auto Scale Manager will try to deploy 'minAsaCount' number of ASAv instances in parallel at one time.</p> <p>STEP: The Auto Scale Manager will deploy the 'minAsaCount' number of ASAv devices one by one at each scheduled interval.</p>	
configurationFile	String	<p>The file path to the ASAv configuration file.</p> <p>Example: https://myserver/asavconfig/asaconfig.txt</p>	N/A
*Azure has restrictions on the naming convention for new resources. Review the limitations or simply use all lowercase. Do not use spaces or any other special characters.			

Auto Scale Deployment

Deploy the Auto Scale ARM Template

The ARM template is used to deploy the resources required by ASAv Auto Scale for Azure. Within a given resource group, the ARM template deployment creates the following:

- Virtual Machine Scale Set (VMSS)
- External Load Balancer
- Internal Load Balancer
- Azure Function App
- Logic App
- Security groups (For Data and Management interfaces)

Before you begin

- Download the ARM template *azure_asav_autoscale.json* from the GitHub repository (<https://github.com/CiscoDevNet/cisco-asav>).

Step 1

If you need to deploy ASAv instances in multiple Azure zones, edit the ARM template based on the zones available in the Deployment region.

Example:

```
"zones": [
  "1",
  "2",
  "3"
],
```

This example shows the “Central US” region which has 3 zones.

Step 2

Edit the traffic rules required in External Load Balancer. You can add any number of rules by extending this ‘json’ array.

Example:

```
{
  "type": "Microsoft.Network/loadBalancers",
  "name": "[variables('elbName')]",
  "location": "[resourceGroup().location]",
  "apiVersion": "2018-06-01",
  "sku": {
    "name": "Standard"
  },
  "dependsOn": [
    "[concat('Microsoft.Network/publicIPAddresses/', variables('elbPublicIpName'))]"
  ],
  "properties": {
    "frontendIPConfigurations": [
      {
        "name": "LoadBalancerFrontEnd",
        "properties": {
          "publicIPAddress": {
            "id": "[resourceId('Microsoft.Network/publicIPAddresses/',
variables('elbPublicIpName'))]"
          }
        }
      }
    ]
  }
}
```

```

    ],
    "backendAddressPools": [
      {
        "name": "backendPool"
      }
    ],
    "loadBalancingRules": [
      {
        "properties": {
          "frontendIPConfiguration": {
            "Id": "[concat(resourceId('Microsoft.Network/loadBalancers', variables('elbName')),
'/frontendIpConfigurations/LoadBalancerFrontend')]"
          },
          "backendAddressPool": {
            "Id": "[concat(resourceId('Microsoft.Network/loadBalancers', variables('elbName')),
'/backendAddressPools/BackendPool')]"
          },
          "probe": {
            "Id": "[concat(resourceId('Microsoft.Network/loadBalancers', variables('elbName')),
'/probes/lbprobe')]"
          },
          "protocol": "TCP",
          "frontendPort": "80",
          "backendPort": "80",
          "idleTimeoutInMinutes": "[variables('idleTimeoutInMinutes')]"
        },
        "Name": "lbrule"
      }
    ]
  ],

```

Note You can also edit this from the Azure portal post-deployment if you prefer not to edit this file.

Step 3

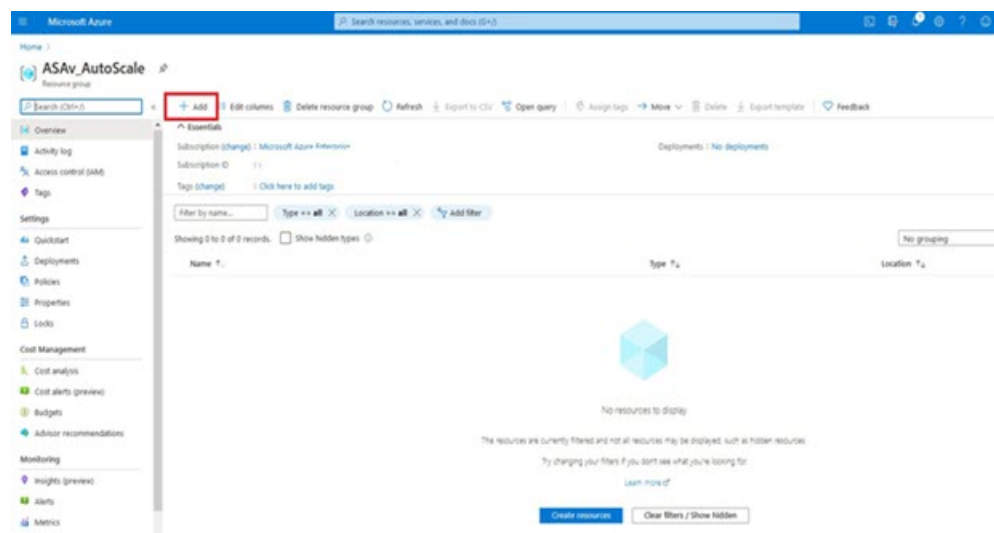
Log in to the Microsoft Azure portal using your Microsoft account username and password.

Step 4

Click **Resource groups** from the menu of services to access the Resource Groups blade. You will see all the resource groups in your subscription listed in the blade.

Create a new resource group or select an existing, empty resource group; for example, *ASAv_AutoScale*.

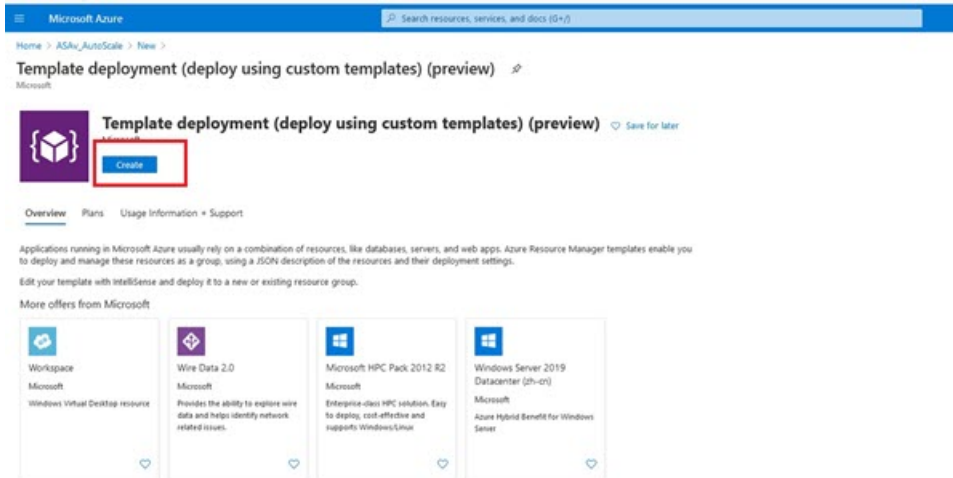
Figure 2: Azure Portal



Step 5 Click **Create a resource (+)** to create a new resource for template deployment. The Create Resource Group blade appears.

Step 6 In **Search the Marketplace**, type **Template deployment (deploy using custom templates)**, and then press **Enter**.

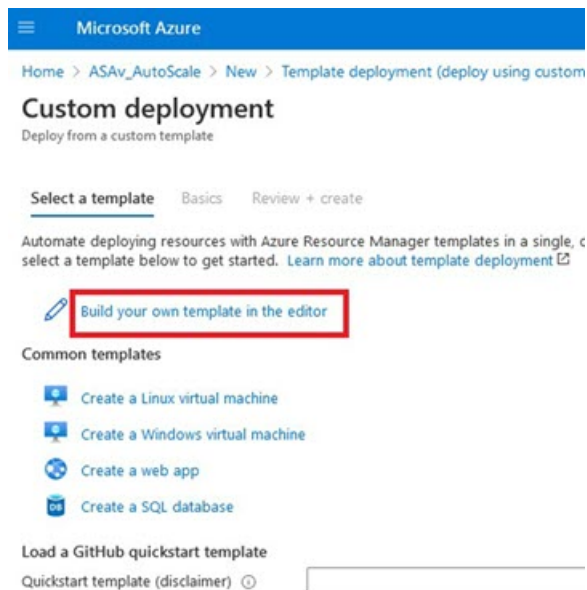
Figure 3: Custom Template Deployment



Step 7 Click **Create**.


Step 8 There are several options for creating a template. Choose **Build your own template in editor**.

Figure 4: Build Your Own Template



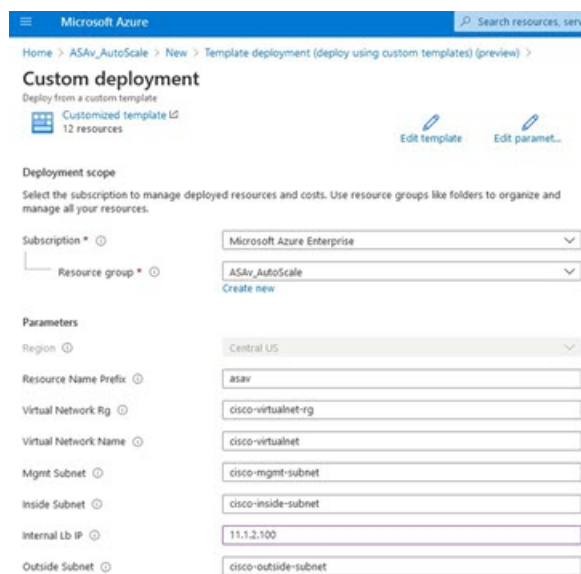
Step 9 In the **Edit template** window, delete all the default content and copy the contents from the updated `azure_asav_autoscale.json` and click **Save**.

Figure 5: Edit Template

**Step 10**

In next section, fill all the parameters. Refer to [Input Parameters, on page 7](#) for details about each parameter, then click **Purchase**.

Figure 6: ARM Template Parameters



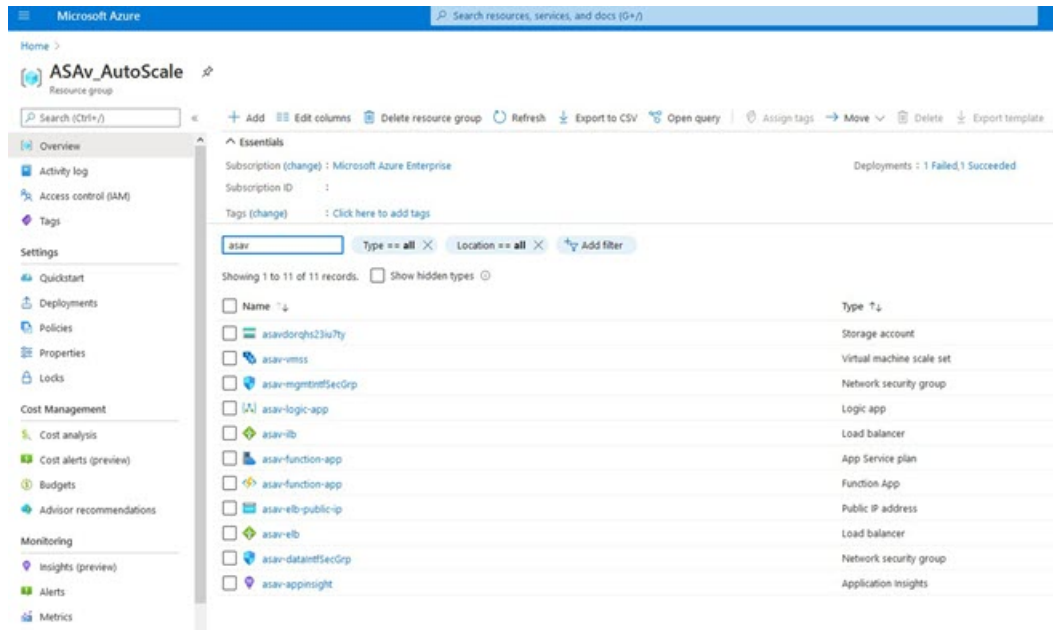
Note You can also click **Edit Parameters** and edit the JSON file or upload pre-filled contents.

The ARM template has limited input validation capabilities, hence it is your responsibility to validate the input.

Step 11

When a template deployment is successful, it creates all the required resources for the ASAv Auto Scale for Azure solution. See the resources in the following figure. The Type column describes each resource, including the Logic App, VMSS, Load Balancers, Public IP address, etc.

Figure 7: ASAv Auto Scale Template Deployment



Deploy the Azure Function App

When you deploy the ARM template, Azure creates a skeleton Function App, which you then need to update and configure manually with the functions required for the Auto Scale Manager logic.

Before you begin

- Build the *ASM_Function.zip* package. See [Appendix - Build Azure Functions from Source Code, on page 27](#).

Step 1 Go to the Function App you created when you deployed the ARM template, and verify that no functions are present. In a browser go to this URL:

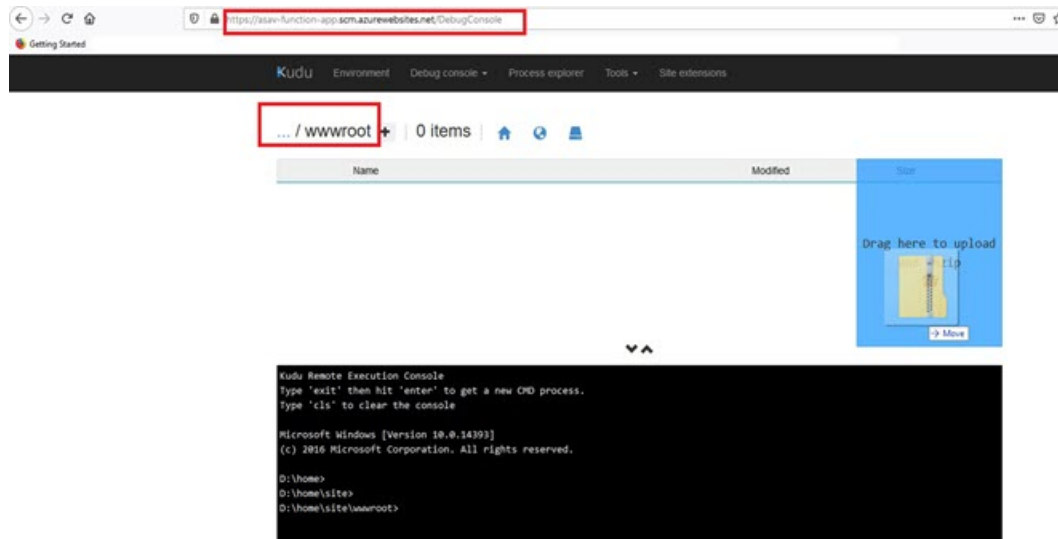
`https://<Function App Name>.scm.azurewebsites.net/DebugConsole`

For the example in [Deploy the Auto Scale ARM Template, on page 10](#):

`https://asav-function-app.scm.azurewebsites.net/DebugConsole`

Step 2 In the file explorer navigate to **site/wwwroot**.

Step 3 Drag-and-drop the **ASM_Function.zip** to the right side corner of the file explorer.

Figure 8: Upload ASAv Auto Scale Functions

Step 4 Once the upload is successful, all of the serverless functions should appear.

Figure 9: ASAv Serverless Functions

The screenshot shows the Azure Portal interface for the ASAv Function App. The address bar indicates the URL: `https://asav-function-app.scm.azurewebsites.net/DebugConsole`. The file explorer shows the `wwwroot` directory with 13 items. A table lists the serverless functions.

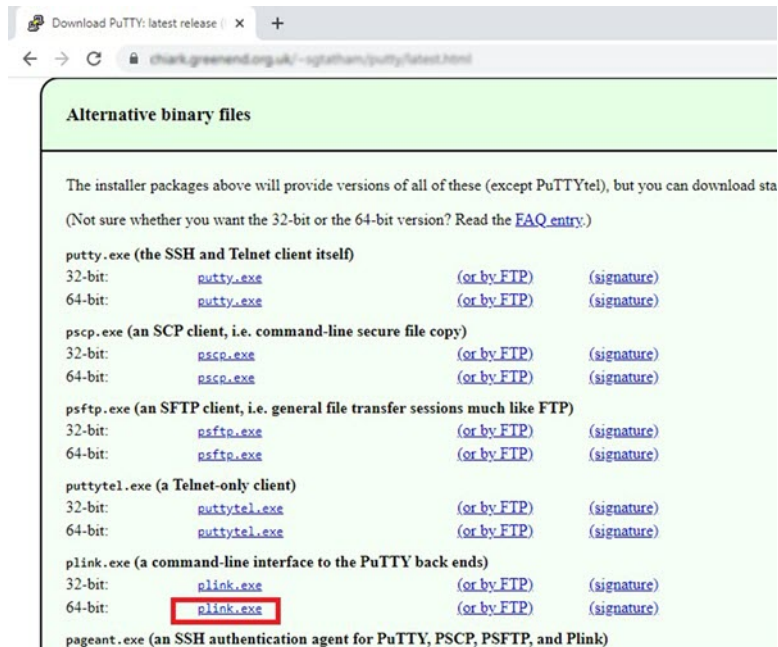
Name	Modified
AsaScaleIn	10/23/2020, 12:28:15 PM
AsaScaleOut	10/23/2020, 12:28:15 PM
AutoScaleManager	10/23/2020, 12:28:16 PM
bin	10/23/2020, 12:28:16 PM
CheckASAvLicenseConfig	10/23/2020, 12:28:27 PM
CleanupASAvConfiguration	10/23/2020, 12:28:27 PM
ConfigureASAv	10/23/2020, 12:28:27 PM

Step 5 Download the PuTTY SSH client.

Azure functions need to access the ASAv via an SSH connection. However, the opensource libraries used in the serverless code do not support the SSH key exchange algorithms used by ASAv. Hence you need to download a pre-built SSH client.

Download the PuTTY command-line interface to the PuTTY back end (*plink.exe*) from www.putty.org.

Figure 10: Download PuTTY



- Step 6** Rename the SSH client executable file **plink.exe** to **asassh.exe**.
- Step 7** Drag-and-drop the **asassh.exe** to the right side corner of the file explorer, to the location where **ASM_Function.zip** was uploaded in the previous step.
- Step 8** Verify the SSH client is present with the function application. Refresh the page if necessary.

Fine Tune the Configuration

There are a few configurations available to fine tune the Auto Scale Manager or to use in debugging. These options are not exposed in the ARM template, but you can edit them under the Function App.

Before you begin

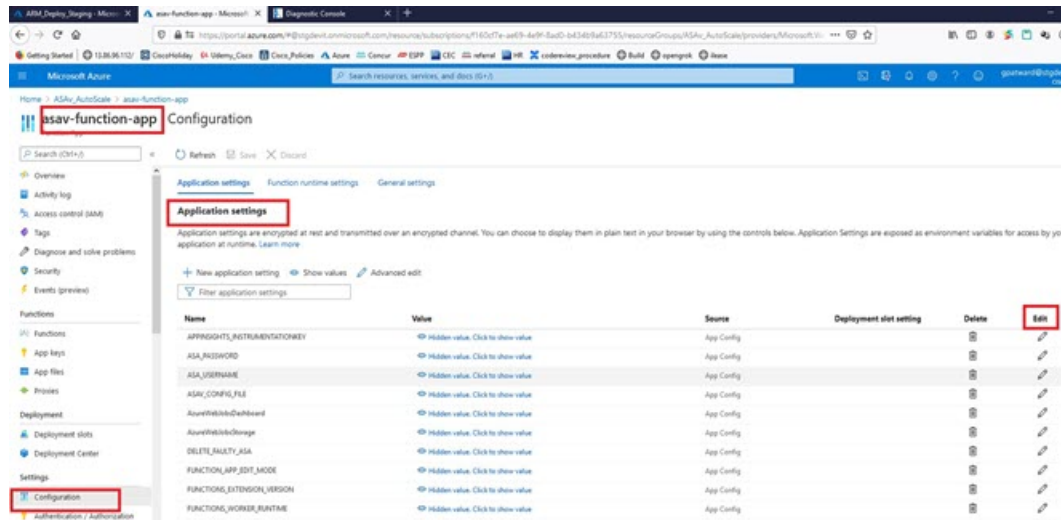


Note This can be edited at any time. Follow this sequence to edit the configurations.

- Disable the Function App.
- Wait for existing scheduled task to finish.
- Edit and save the configuration.
- Enable the Function App.

- Step 1** In the Azure portal, search for and select the ASAv function application.

Figure 11: ASAv Function Application

**Step 2**

Configurations passed via the ARM template can also be edited here. Variable names may appear different from the ARM template, but you can easily identify the purpose of these variables from their name.

Most of the options are self-explanatory from the name. For example:

- Configuration Name: “DELETE_FAULTY_ASA” (Default value : YES)

During Scale-Out a new ASAv instance is launched and configured via the configuration file. In case the configuration fails, based on this option Auto Scale Manager will decide to keep that ASAv instance or delete it. (YES : Delete faulty ASAv / NO : Keep the ASAv instance even if the configuration fails).

- In the Function App settings, all the variables (including variables containing a secure string like ‘password’) can be seen in clear text format by users that have access to the Azure subscription.

If users have any security concerns with this (for example, if an Azure subscription is shared among users with lower privileges within the organization), a user can make use of Azure’s *Key Vault* service to protect passwords. Once this is configured, instead of providing a clear text ‘password’ in function settings, a user has to provide a secure identifier generated by the key vault where the password is stored.

Note Search the Azure documentation to find the best practices to secure your application data.

Configure the IAM Role in the Virtual Machine Scale Set

Azure Identity and Access Management (IAM) is used as a part of Azure Security and Access Control to manage and control a user’s identity. Managed identities for Azure resources provides Azure services with an automatically managed identity in Azure Active Directory.

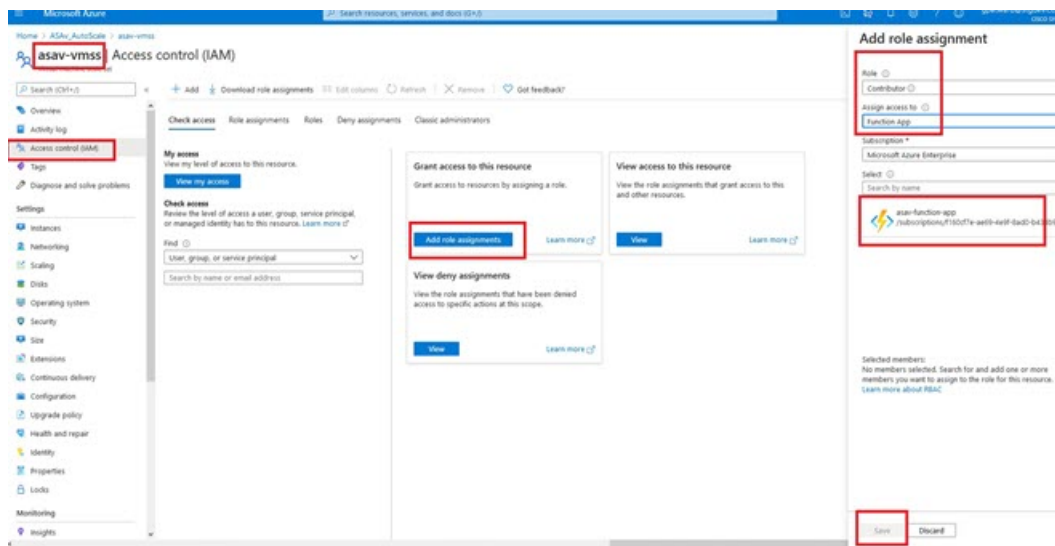
This allows the Function App to control the Virtual Machine Scale Sets (VMSS) without explicit authentication credentials.

Step 1

In the Azure portal, go to the VMSS.

- Step 2** Click **Access control (IAM)**.
- Step 3** Click **Add** to add a role assignment
- Step 4** From the **Add role assignment** drop-down, choose **Contributor**.
- Step 5** From the **Assign access to** drop-down, choose **Function App**.
- Step 6** Select the ASAv function application.

Figure 12: AIM Role Assignment



- Step 7** Click **Save**.

Note You should also verify that there are no ASAv instances launched yet.

Update Security Groups

The ARM template creates two security groups, one for the Management interface, and one for data interfaces. The Management security group will allow only traffic required for ASAv management activities. However, the data interface security group will allow all traffic.

Fine tune the security group rules based on the topology and application needs of your deployments.

Note The data interface security group should allow at a minimum SSH traffic from the load balancers.

Update the Azure Logic App

The Logic App acts as the orchestrator for the Autoscale functionality. The ARM template creates a skeleton Logic App, which you then need to be update manually to provide the information necessary to function as the Auto Scale orchestrator.

Step 1

From the repository, retrieve the file *LogicApp.txt* to the local system and edit as shown below.

Important Read and understand all of these steps before proceeding.

These manual steps are not automated in the ARM template so that only the Logic App can be upgraded independently later in time.

- Required: Find and replace all the occurrences of “SUBSCRIPTION_ID” with your subscription ID information.
- Required: Find and replace all the occurrences of “RG_NAME” with your resource group name.
- Required: Find and replace all of the occurrences of “FUNCTIONAPPNAME” to your function app name. The following example shows a few of these lines in the *LogicApp.txt* file:

```
"AutoScaleManager": {
  "inputs": {
    "function": {
      "id":
"/subscriptions/SUBSCRIPTION_ID/resourceGroups/RG_NAME/providers/Microsoft.Web/sites/FUNCTIONAPPNAME/functions/AutoScaleManager"
    }
  },
  "Deploy_Changes_to_ASA": {
    "inputs": {
      "body": "@body('AutoScaleManager')",
      "function": {
        "id":
"/subscriptions/SUBSCRIPTION_ID/resourceGroups/RG_NAME/providers/Microsoft.Web/sites/FUNCTIONAPPNAME/functions/DeployConfiguration"
      }
    },
    "DeviceDeRegister": {
      "inputs": {
        "body": "@body('AutoScaleManager')",
        "function": {
          "id":
"/subscriptions/SUBSCRIPTION_ID/resourceGroups/RG_NAME/providers/Microsoft.Web/sites/FUNCTIONAPPNAME/functions/DeviceDeRegister"
        }
      },
      "runAfter": {
        "Delay_For_connection_Draining": [

```

- (Optional) Edit the trigger interval, or leave the default value (5). This is the time interval at which the Autoscale functionality is periodically triggered. The following example shows these lines in the *LogicApp.txt* file:

```
"triggers": {
  "Recurrence": {
    "conditions": [],
    "inputs": {},
    "recurrence": {
      "frequency": "Minute",
      "interval": 5
    }
  },

```

- e) (Optional) Edit the time to drain, or leave the default value (5). This is the time interval to drain existing connections from the ASAv before deleting the device during the Scale-In operation. The following example shows these lines in the *LogicApp.txt* file:

```
"actions": {
  "Branch_based_on_Scale-In_or_Scale-Out_condition": {
    "actions": {
      "Delay_For_connection_Draining": {
        "inputs": {
          "interval": {
            "count": 5,
            "unit": "Minute"
          }
        }
      }
    }
  }
}
```

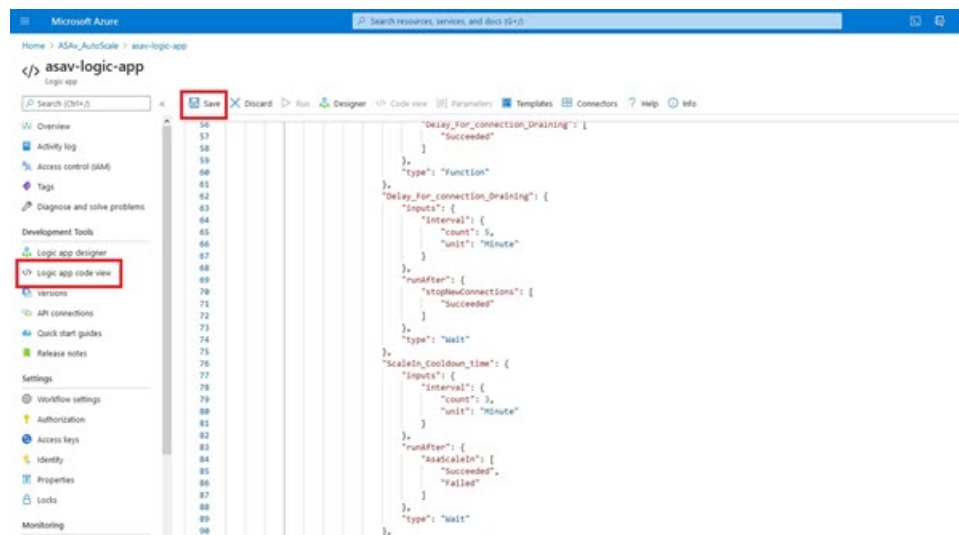
- f) (Optional) Edit the cool down time, or leave the default value (10). This is the time to perform NO ACTION after the Scale-Out is complete. The following example shows these lines in the *LogicApp.txt* file:

```
"actions": {
  "Branch_based_on_Scale-Out_or_Invalid_condition": {
    "actions": {
      "Cooldown_time": {
        "inputs": {
          "interval": {
            "count": 10,
            "unit": "Second"
          }
        }
      }
    }
  }
}
```

Note These steps can also be done from the Azure portal. Consult the Azure documentation for more information.

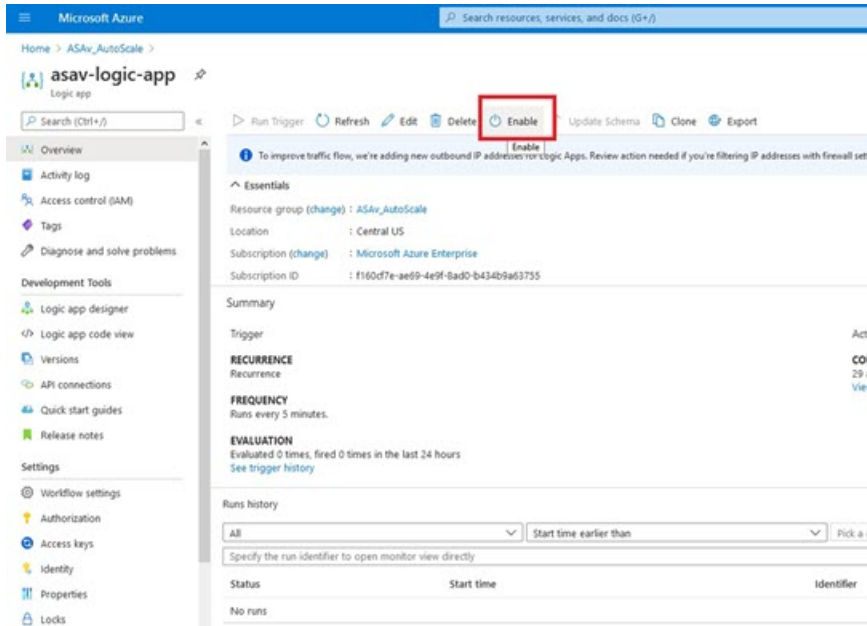
Step 2 Go to the **Logic App code view**, delete the default contents and paste the contents from the edited *LogicApp.txt* file, and click **Save**.

Figure 13: Logic App Code View



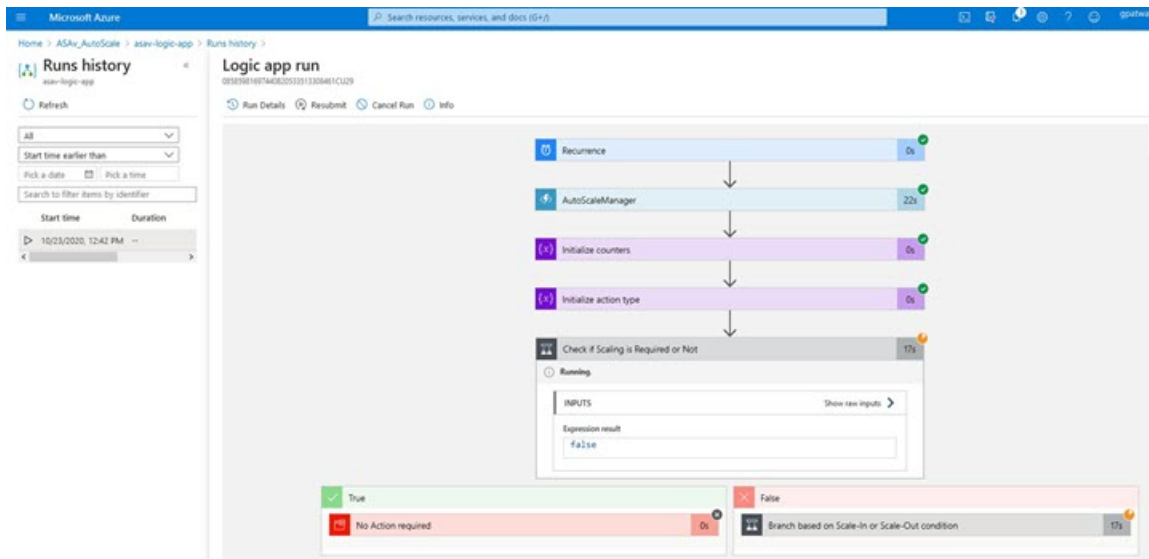
Step 3 When you save the Logic App it is in a 'Disabled' state. Click **Enable** when you want to start the Auto Scale Manager.

Figure 14: Enable Logic App



Step 4 Once enabled the tasks start running. Click the 'Running' status to see the activity.

Figure 15: Logic App Running Status



Step 5 Once the Logic App starts, all of the the deployment-related steps are complete.

Step 6 Verify in the VMSS that ASAv instances are being created.

Figure 16: ASAv Instances Running

Name	Computer name	Status	Health state	Provisioning
asav-vmss_0	asav-vmss000000	Creating (Running)		Creating
asav-vmss_1	asav-vmss000001	Creating (Running)		Creating
asav-vmss_2	asav-vmss000002	Creating (Running)		Creating

In this example, three ASAv instances are launched because 'minAsaCount' was set to '3' and 'initDeploymentMode' was set to 'BULK' in the ARM template deployment.

Upgrade the ASAv

The ASAv upgrade is supported only in the form of an image upgrade of virtual machine scale set (VMSS). Hence you upgrade the ASAv through the Azure REST API interface.



Note You can use any REST client to do upgrade the ASAv. The following is a simple example.

Before you begin

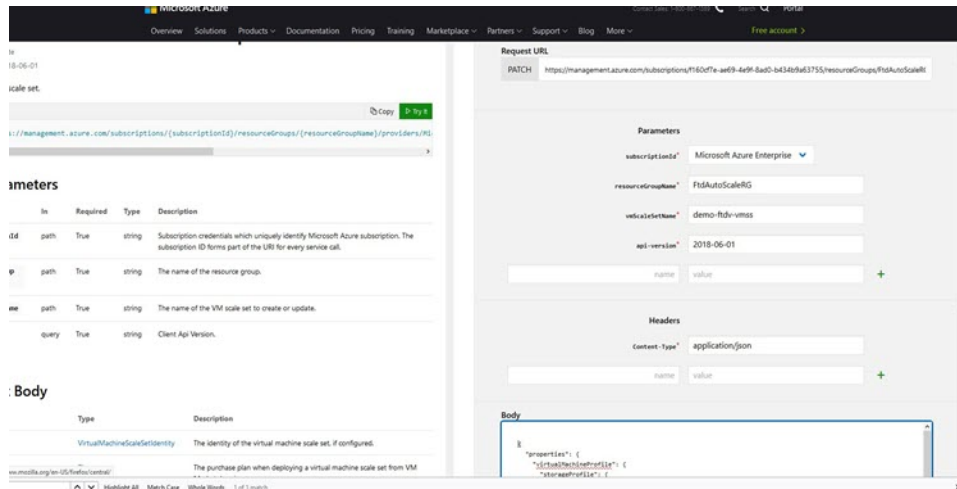
- Obtain the new ASAv image version available in market place (example: 914.001).
- Obtain the SKU used to deploy original scale set (example: asav-azure-byol).
- Obtain the Resource Group and the virtual machine scale set name

Step 1 In a browser go to the following URL:

<https://docs.microsoft.com/en-us/rest/api/compute/virtualmachinescalesets/update#code-try-0>

Step 2 Enter the details in the parameter section.

Figure 17: Upgrade the ASAv



Step 3 Enter the JSON input containing new the ASAv image version, SKU, and trigger RUN in the **Body** section.

```
{
  "properties": {
    "virtualMachineProfile": {
      "storageProfile": {
        "imageReference": {
          "publisher": "cisco",
          "offer": "cisco-asav",
          "sku": "asav-azure-byol",
          "version": "650.32.0"
        }
      }
    }
  }
}
```

Step 4 A successful response from Azure means that the VMSS has accepted the change.

The new image will be used in the new ASAv instances which will get launched as part of Scale-Out operation.

- Existing ASAv instances will continue to use the old software image while they exist in a scale set.
- You can override the above behavior and upgrade the existing ASAv instances manually. To do this, click the **Upgrade** button in the VMSS. It will reboot and upgrade the selected ASAv instances. You must reregister and reconfigure these upgraded ASAv instances manually. **Note that this method is NOT recommended.**

Auto Scale Logic

Scale-Out Logic

- **POLICY-1:** Scale-Out will be triggered when the average load of **any** ASAv goes beyond the Scale-Out threshold for the configured duration.
- **POLICY-2:** Scale-Out will be triggered when average load of **all** of the ASAv devices go beyond Scale-Out threshold for the configured duration.

Scale-In Logic

- If the CPU utilization of **all** of the ASAv devices goes below the configured Scale-In threshold for the configured duration.

Notes

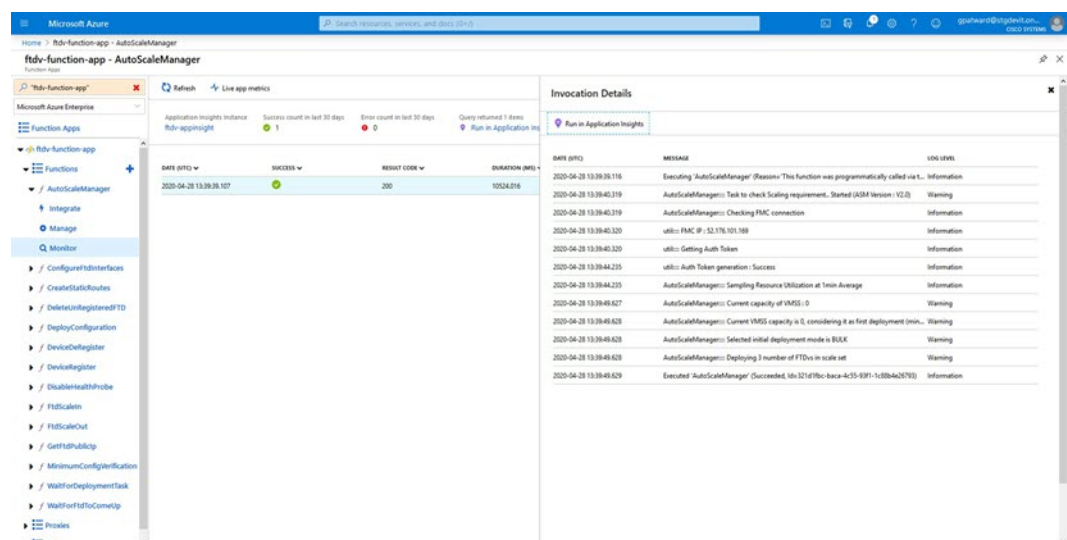
- Scale-In/Scale-Out occurs in steps of 1 (i.e. only 1 ASAv will be scaled in/out at a time).
- The above logic is based on the assumption that the load balancer will try to equally distribute connections across all ASAv devices, and on an average all ASAv devices should be loaded equally.

Auto Scale Logging and Debugging

Each component of the serverless code has its own logging mechanism. In addition, logs are published to application insight.

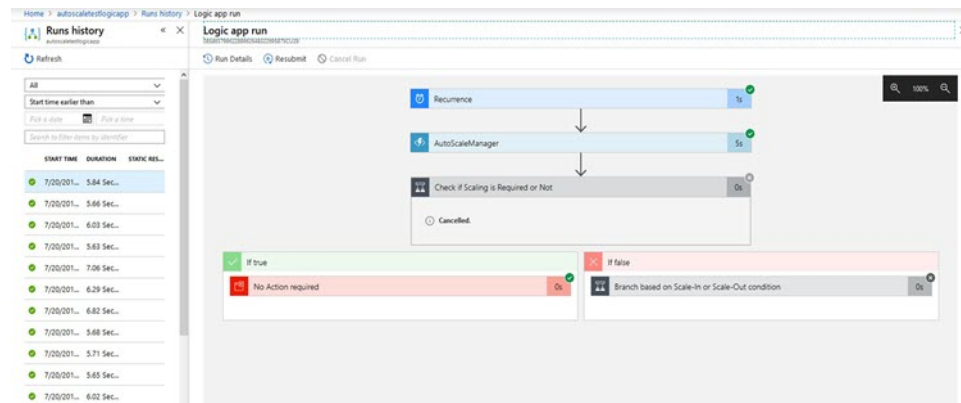
- Logs of individual Azure functions can be viewed.

Figure 18: Azure Function Logs



- Similar logs for each run of the Logic App and its individual components can be viewed.

Figure 19: Logic App Run Logs



- If needed, any running task in the Logic App can be stopped/terminated at any time. However, currently running ASAv devices getting launched/terminated will be in an inconsistent state.
- The time taken for each run/individual task can be seen in the Logic App.
- The Function App can be upgraded at any time by uploading a new zip. Stop the Logic App and wait for all tasks to complete before upgrading the Function App.

Auto Scale Guidelines and Limitations

Be aware of the following guidelines and limitations when deploying ASAv Auto Scale for Azure:

- Scaling decisions are based on CPU utilization.
- The ASAv Management interface is configured to have public IP address.
- Only IPv4 is supported.
- The ARM template has limited input validation capabilities, hence it is your responsibility to provide the correct input validation.
- The Azure administrator can see sensitive data (such as admin login credentials and passwords) in plain text format inside Function App environment. You can use the *Azure Key Vault* service to secure sensitive data.

Auto Scale Troubleshooting

The following are common error scenarios and debugging tips for ASAv Auto Scale for Azure:

- Unable to SSH into the ASAv: Check if a complex password is passed to the ASAv via the template; check if Security Groups allow SSH connections.
- Load Balancer Health check failure: Check if the ASAv responds to SSH on data interfaces; check Security Group settings.

- Traffic issues: Check Load Balancer rules, NAT rules / Static routes configured in ASAv; check Azure virtual network / subnets / gateway details provided in the template and Security Group rules.
- Logic App failed to access VMSS: Check if the IAM role configuration in VMSS is correct.
- Logic App runs for very long time: Check SSH access on scaled-out ASAv devices; check the state of the ASAv devices in Azure VMSS.
- Azure Function throwing error related to subscription ID : Verify that you have a default subscription selected in your account.
- Failure of Scale-In operation: Sometimes Azure takes considerably long time to delete an instance, in such situation Scale-in operation may time out and report error but eventually the instance will get deleted.
- Before doing any configuration change, make sure to disable the logic application and wait for all the running tasks to complete.

Appendix - Build Azure Functions from Source Code

System Requirements

- Microsoft Windows desktop/laptop.
- Visual Studio (tested with Visual studio 2019 version 16.1.3)



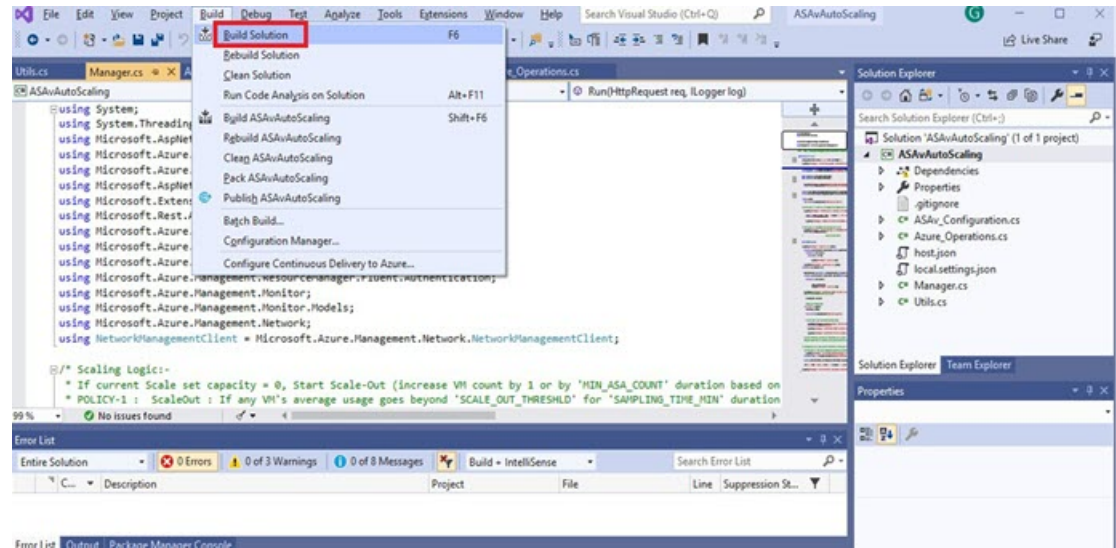
Note Azure functions are written using C#.

- The "Azure Development" workload needs to be installed in Visual Studio.

Build with Visual Studio

1. Download the 'code' folder to the local machine.
2. Navigate to the folder 'ASAAutoScaling'.
3. Open the project file 'ASAAutoScaling.csproj' in Visual Studio.
4. Use Visual Studio standard procedure to Clean and Build.

Figure 20: Visual Studio Build



5. Once the build is compiled successfully, navigate to the `\bin\Release\netcoreapp2.1` folder.
6. Select all the contents, click **Send to > Compressed (zipped) folder**, and save the ZIP file as `ASM_Function.zip`.

Figure 21: Build ASM_Function.zip

